




DEEP DIVE INTO KUBERNETES

REAL-WORLD USE CASES & INTERNAL FLOWS




1 How Kubernetes Actually Decides Where a Pod Runs

↳ Scheduler deep dive: filtering, scoring, and node selection explained




2 What Really Happens When You Run `kubectl apply -f`

↳ End-to-end flow from API server → etcd → controllers → reconciliation




3 Kubernetes Deployment Rollout Strategies Explained

↳ Rolling updates vs Recreate vs Blue-Green deployments in production




4 Why Your Pod Keeps Restarting (CrashLoopBackOff Explained)

↳ Common root causes + how to debug like a pro




5 Kubernetes Service Discovery Explained

↳ How Pods actually find and communicate with each other




6 Kubernetes Networking Under the Hood

↳ CNI plugins, kube-proxy, and packet routing flow explained




7 ImagePullBackOff vs ErrImagePull

↳ Real differences, causes, and practical fixes




8 How Kubernetes Autoscaling Really Works

↳ HPA, VPA, and Cluster Autoscaler in action




9 What Happens When a Node Goes NotReady

↳ Failure detection, rescheduling, and self-healing behavior




10 ConfigMap vs Secret in Production

↳ When and how to use them correctly in real systems



11 How Kubernetes RBAC Works Internally

↳ From request authentication → authorization decision flow



12 End-to-End Pod Lifecycle Explained

↳ From creation → scheduling → running → termination

